


Writing Portable Building Analytics with the Brick Metadata Schema

Gabe Fierro, UC Berkeley

ACM E-Energy

gtfierro@eecs.berkeley.edu



Outline

1. Lecture on Brick Schema
2. Overview of Mortar platform
3. Guided exploration of Mortar with online tutorial

Resources

Brick Website

<https://brickschema.org/>

Mortar Website

<https://mortardata.org/>

Tutorial Page

<https://tutorial.mortardata.org/>

Brick Explore Tool

<https://querybuilder.mortardata.org/>

Applications

Demand Response

Occupant
Interaction

NILM

Occupancy Models

Predictive Control

Fault Detection

Management Services

APIs

Data Storage

Access Control

Monitoring

Search

Privacy

Buildings

Residential

Large Commercial

Factory

Research Lab

Small Commercial

Hospital

Sensors Equipment

HVAC

Appliances

Lighting

Fire Safety

Conditioning

Metering

Heterogeneity

- Different BMS, equipment vendors
- Custom-designed controls, systems, architecture
- All of this changes over time

Buildings

Residential

Large Commercial

Factory

Research Lab

Small Commercial

Hospital

Sensors Equipment

HVAC

Appliances

Lighting

Fire Safety

Conditioning

Metering

Existing Metadata Schemata

- Survey 90 apps from building science literature
- Decompose apps into “things” and “relationships”
- Can existing schemas describe what’s needed?

- **NO!**

	IFC	SSN	Haystack
Tag Coverage	29%	11%	54%
Relationships	n/a	Only spatial	n/a

Bhattacharya, Arka, Joern Ploennigs, and David Culler. "Short Paper: Analyzing Metadata Schemas for Buildings: The Good, the Bad, and the Ugly." *BuildSys*, 2015.

Applications

Demand Response

Occupant
Interaction

NILM

Occupancy Models

Predictive Control

Fault Detection

Management Services

APIs

Data Storage

Access Control

Monitoring

Search


Privacy

BRICK Metadata Schema

- Write applications against BRICK
- Apps become *agnostic* to the underlying hardware
- No more hard-coding points in applications
- “Write once, run anywhere”

Brick Schema

- Graph-based metadata schema for smart buildings
- Capture physical, logical, virtual **entities** in buildings using a **class hierarchy**
- Capture the necessary **relationships** between them
- Use Brick to describe timeseries data and its context



Brick

<https://brickschema.org>

Berkeley
UNIVERSITY OF CALIFORNIA

Carnegie
Mellon
University

Johnson
Controls



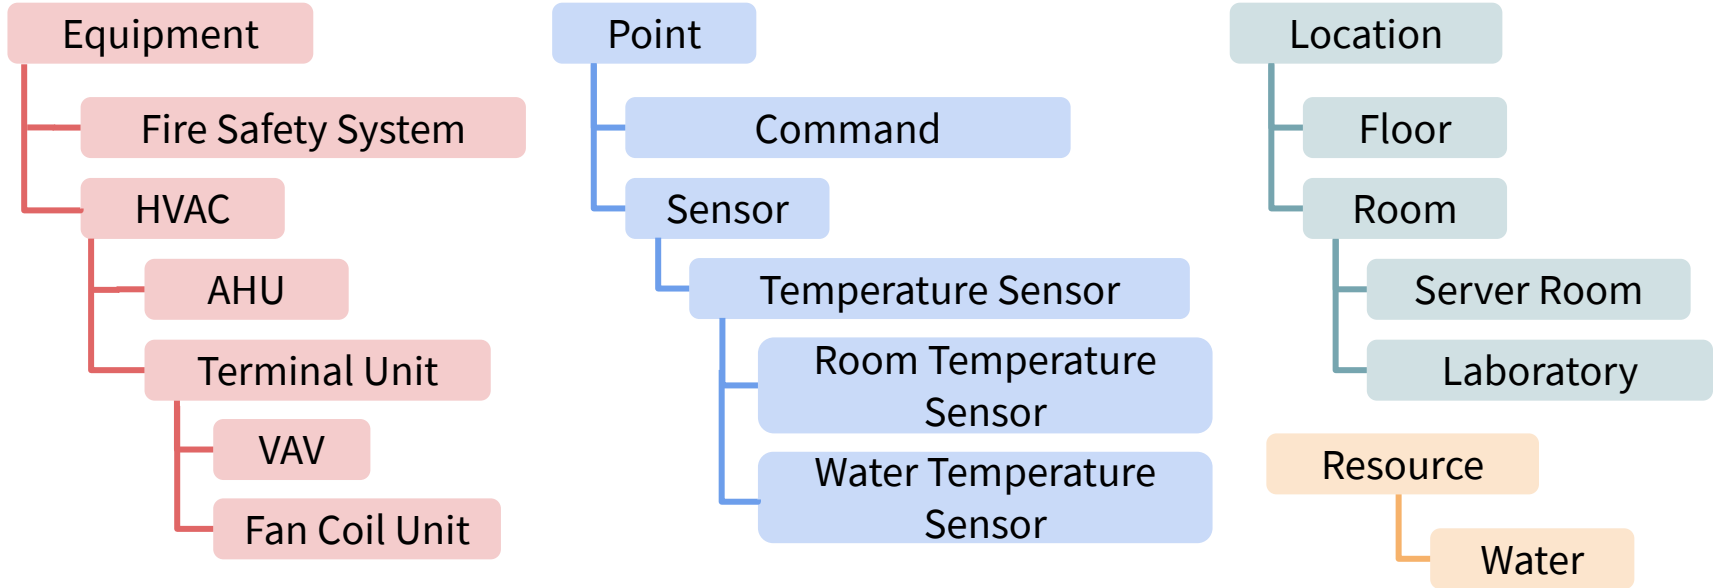
UC San Diego



UCLA SDU

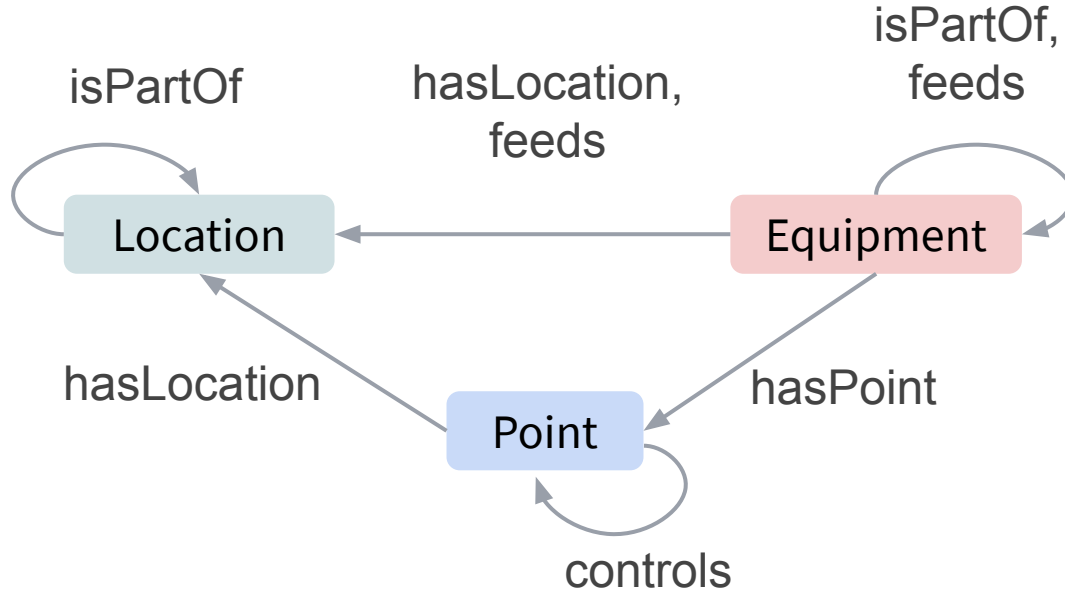


Brick Schema: Class Structure



- Standardized class structure enables discoverability
- Extensible: allow site/deployment-specific classes

Brick Schema: Relationships



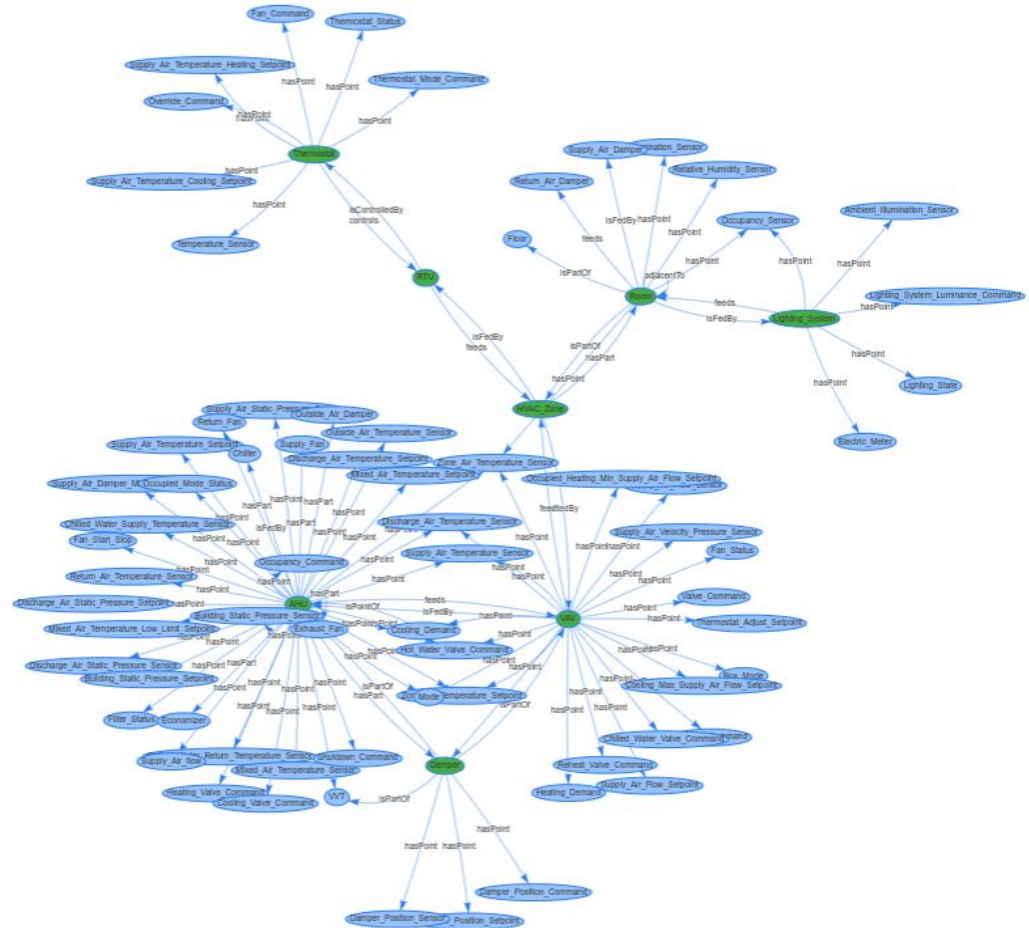
- Relationships capture how “things” relate to each other
- Relationships can be transitive, symmetrical
- Help Brick extend to cover new settings, equipment

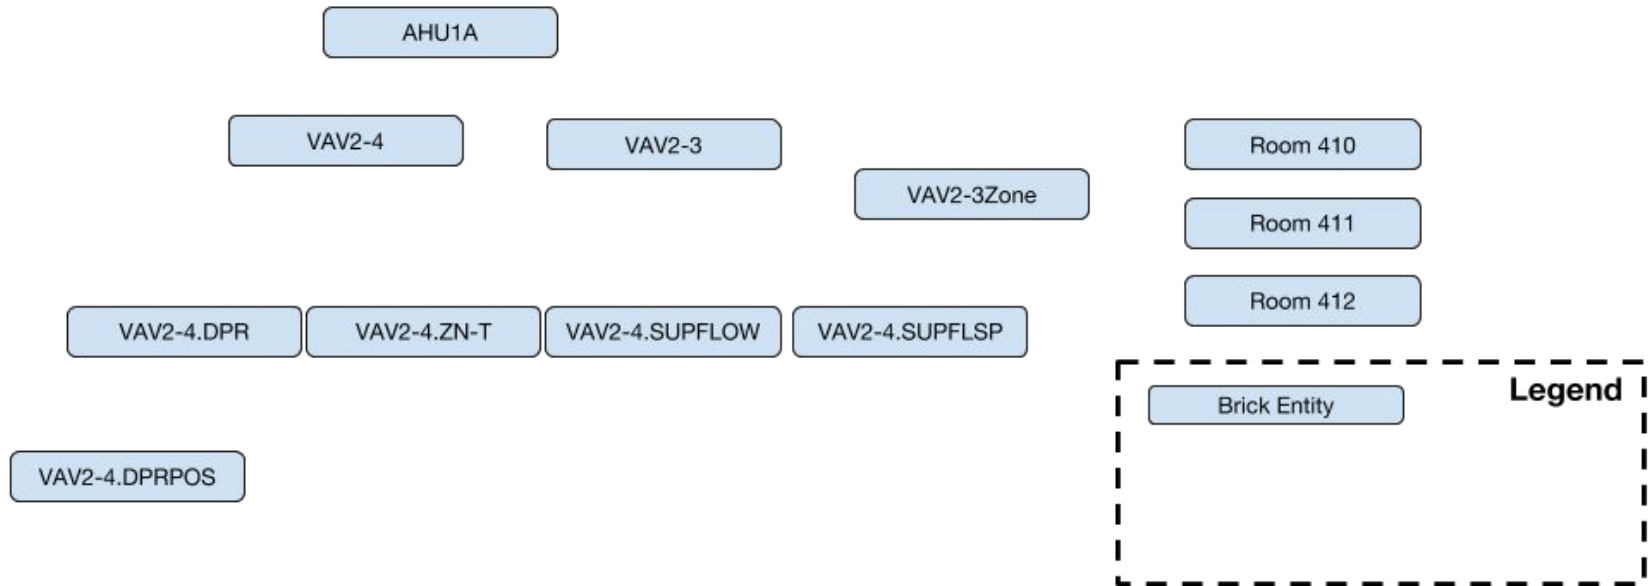
Brick Schema: Relationships

Relationship	Definition
contains/isLocatedIn	Spatial location of the subject: room, building, campus, city chain
controls/isControlledBy	Identifies the affected subject of some control block (probably a functional block)
hasPart/isPartOf	Mechanical composition
hasPoint/isPointOf	Associates measurement points/timeseries with the entity it measures for
feeds/isFedBy	The passage of some medium: light, air, power, water, etc
type/isTypeOf	Instantiates “things”

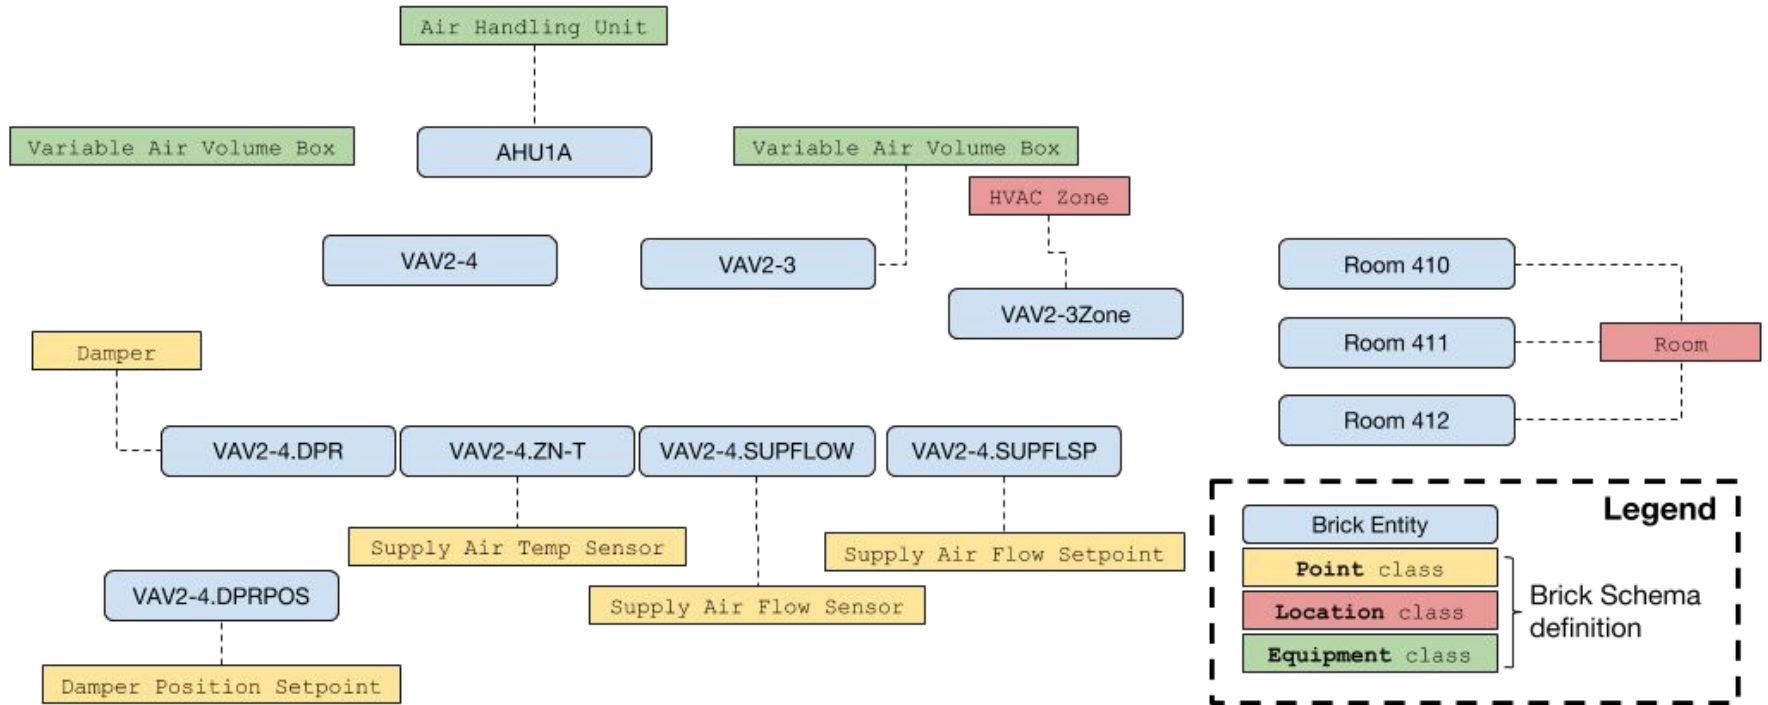
A Brick Model is a Graph

- **Nodes = “things”**
 - Building assets
 - Equipment
 - Subsystems
 - Class structure
- **Edges = “relationships”**
 - Location
 - Control
 - Connectivity
 - Composition
 - etc

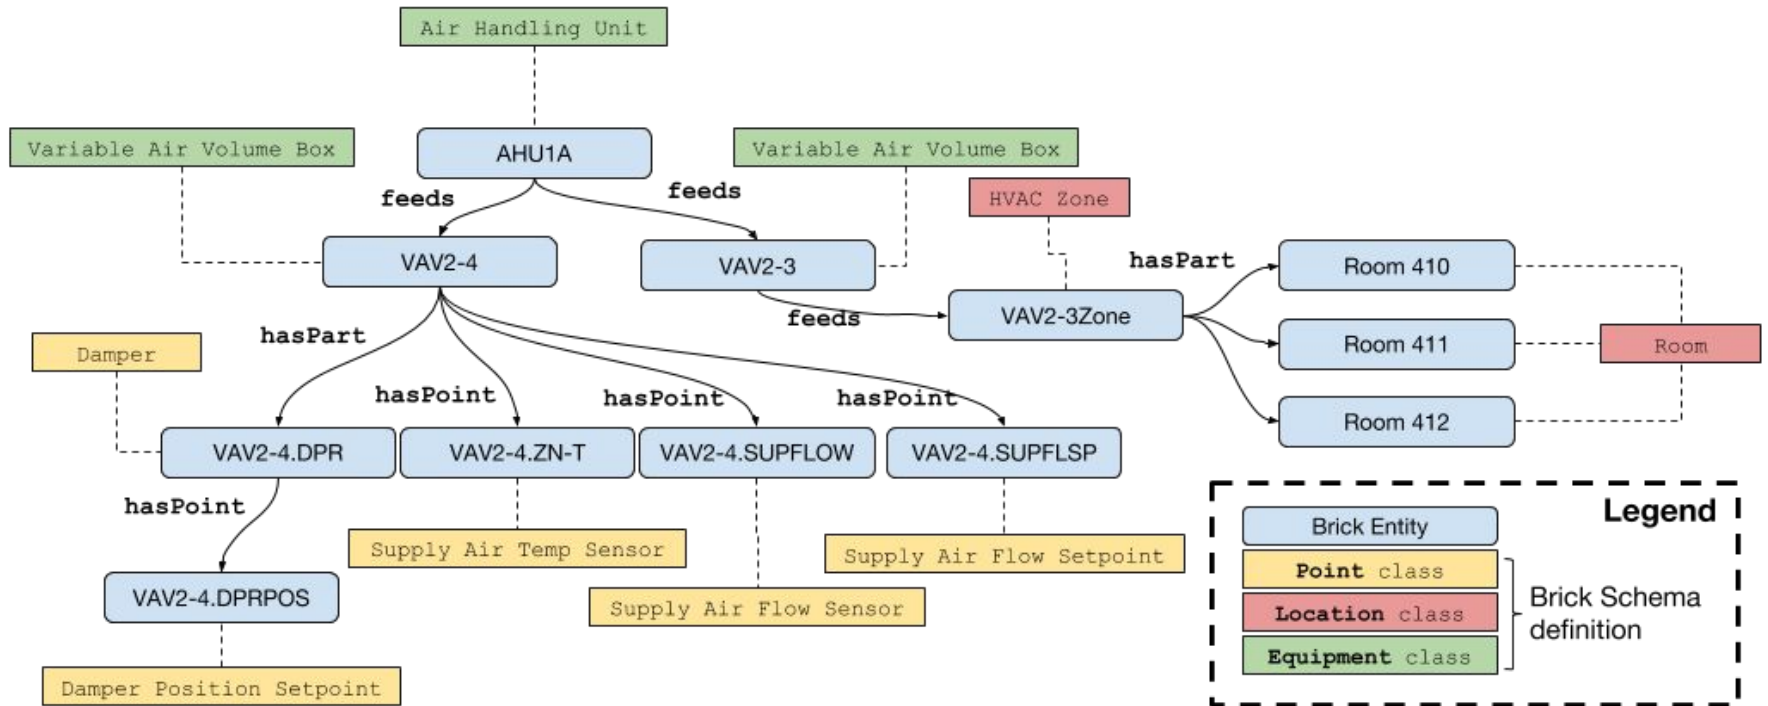





- Have a set of physical, virtual “*things*” and points that an application wants to refer to



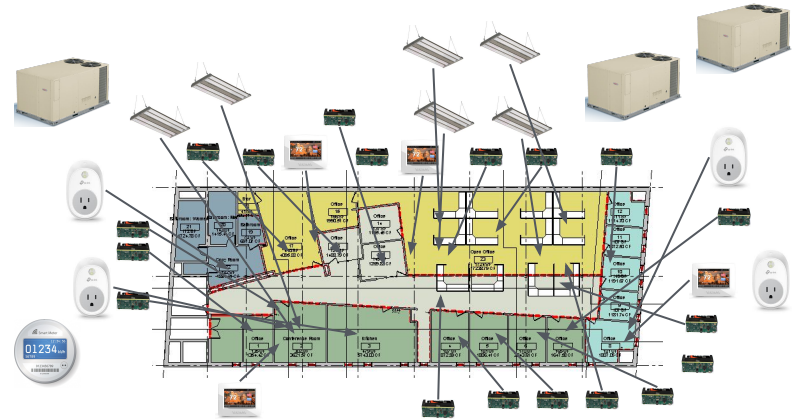
- Have a set of physical, virtual “*things*” and points that an application wants to refer to
- Brick defines a **hierarchical class structure** to define **standard names** for equipment, points, locations, etc



- Have a set of physical, virtual “*things*” and points that an application wants to refer to
- Brick defines a **hierarchical class structure** to define **standard names** for equipment, points, locations, etc
- Brick defines a set of **standard relationships** that describe how things are **connected**



represents



A Brick model represents the assets and relationships and data in a building

Query

```

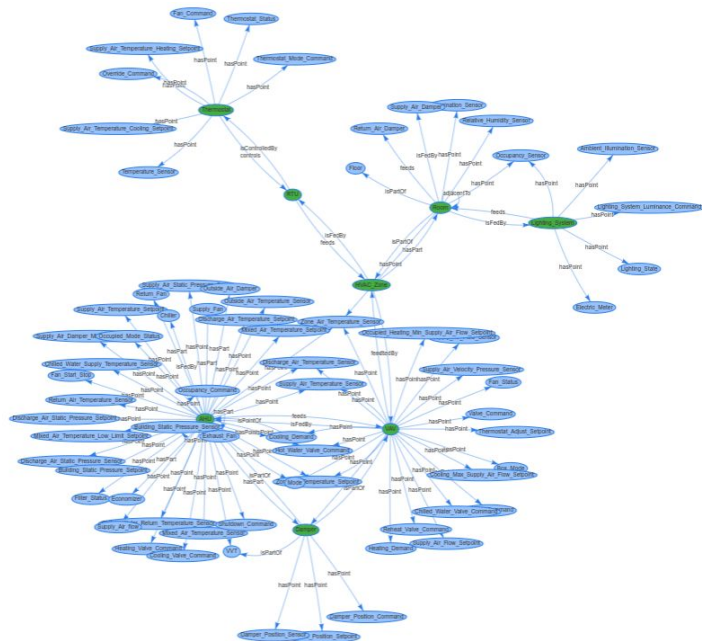
1 SELECT ?tstat ?zone ?tstat_uri ?state_uuid ?temp_uuid ?weather_uuid FROM ciece WHERE {
2   ?zone rdf:type brick:HVAC_Zone .
3   ?rtu bf:feeds+ ?zone .
4
5   ?tstat rdf:type/rdfs:subClassOf* brick:Thermostat .
6   ?tstat bf:controls+ ?rtu .
7   ?tstat bf:uri ?tstat_uri .
8
9   ?tstat bf:hasPoint ?state .
10  ?state rdf:type brick:Thermostat_HVAC_Operation_Status .
11  ?state bf:uuid ?state_uuid .
12
13  ?tstat bf:hasPoint ?temp .
14  ?temp rdf:type/rdfs:subClassOf* brick:Temperature_Sensor .
15  ?temp bf:uuid ?temp_uuid .
16
17  ?tstat bf:hasSite ?site .
18  ?weather_sensor bf:hasSite ?site .
19  ?weather_sensor rdf:type/rdfs:subClassOf* brick:Temperature_Sensor .
20  ?weather_sensor bf:uuid ?weather_uuid
21 };

```



Application

queries






Brick model

An application **queries** a Brick model to retrieve the data + configuration it needs

Query

```
1 SELECT ?tstat ?zone ?tstat_uri ?state_uuid ?temp_uuid ?weather_uuid FROM ciee WHERE {  
2   ?zone rdf:type brick:HVAC_Zone .  
3   ?rtu bf:feeds+ ?zone .  
4  
5   ?tstat rdf:type/rdfs:subClassOf* brick:Thermostat .  
6   ?tstat bf:controls+ ?rtu .  
7   ?tstat bf:uri ?tstat_uri .  
8  
9   ?tstat bf:hasPoint ?state .  
10  ?state rdf:type brick:Thermostat_HVAC_Operation_Status .  
11  ?state bf:uuid ?state_uuid .  
12  
13  ?tstat bf:hasPoint ?temp .  
14  ?temp rdf:type/rdfs:subClassOf* brick:Temperature_Sensor .  
15  ?temp bf:uuid ?temp_uuid .  
16  
17  ?tstat bf:hasSite ?site .  
18  ?weather_sensor bf:hasSite ?site .  
19  ?weather_sensor rdf:type/rdfs:subClassOf* brick:Temperature_Sensor .  
20  ?weather_sensor bf:uuid ?weather_uuid .  
21 };
```

executes



Queries allow apps to account for building heterogeneity and **customize their operation** to each building.

This is called **application portability**

Outline

- ~~1. Lecture on Brick Schema~~
2. Overview of Mortar platform
3. Guided exploration of Mortar with online tutorial

Resources

Brick Website

<https://brickschema.org/>

Mortar Website

<https://mortardata.org/>

Tutorial Page

<https://tutorial.mortardata.org/>

Brick Explore Tool

<https://querybuilder.mortardata.org/>

How do we build a platform that enables
reproducible building research?



Why are Portable Applications Hard to Write?



- Hardcoded point names:
 - *Lights in this building are* WS86004.RELAYXX
- Assumptions about building subsystems
 - *This building has room-level temperature sensors*
- Tightly-coupled phases of operation:
 - Downloading, cleaning data often specific to the building at hand
- Access to different buildings:
 - Why write portably if you don't have to?
- How to name, describe and discover data
 - **Brick** gives us this abstraction: describe data in terms of types and relationships

Mortar Platform Contributions

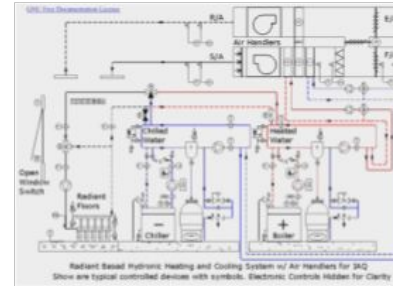

- Open data set of buildings with timeseries data and Brick models
- Modular, extensible architecture for portable analytics applications
- Data platform for the development, execution and evaluation of portable analytics









Air-based HVAC Systems







Radiant Heating/Cooling Systems

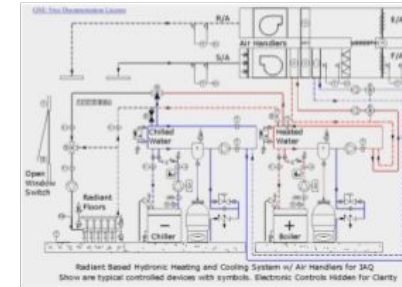


Rooftop Unit Systems






Air-based HVAC Systems



Radiant Heating/Cooling Systems



Rooftop Unit Systems



Indoor
Temperature
Prediction
(Building 1)

Custom
Implementations



Indoor
Temperature
Prediction
(Building 2)



Indoor
Temperature
Prediction
(Building 3)

Indoor
Temperature
Prediction
(Building 4)

Indoor
Temperature
Prediction
(Building 5)

Facilitating reproducibility means reducing the effort of implementation/evaluation across contexts:


1. Expressive description of data + context (Brick)
2. Large, diverse dataset
3. Application-specific normalization of data



Systems



ing Systems



Rooftop Unit Systems



Open Dataset

- Live data for > 100 buildings
 - Up to 6 years of data
- Contextualized timeseries data
 - Each building has a Brick model
- Public API for data access
 - **pymortar** Python library
- Open source library of analytics applications

SoftwareDefinedBuildings / [mortar-analytics](#) Watch 489 Star 19,384 Fork 34

[Code](#) [Issues 3](#) [Pull requests 6](#) [Projects 0](#) [Insights](#)

Repository for implementations of analytics for mortardata.org

97 commits 1 branch 0 releases 3 contributors BSD-2-Clause

https://mortardata.org

MortarData

Menu

- Sign Up
- Documentation
- Mortar Analytics Library

Mortar is in Alpha! Please [subscribe](#) to the mortar newsletter

What is Mortar?

Access to large amounts of real-world data has long been a challenge in the building environment. Open data sets exist, but they are limited in scope and how it is described).

The goal of Mortar is to provide a large, diverse and accessible evaluation of building analytics.

At this time, Mortar contains **107 buildings**, spanning a wide range of building types, provided by a [Brick model](#). The Brick model describes how the building equipment exists and how it is monitored, (3) the

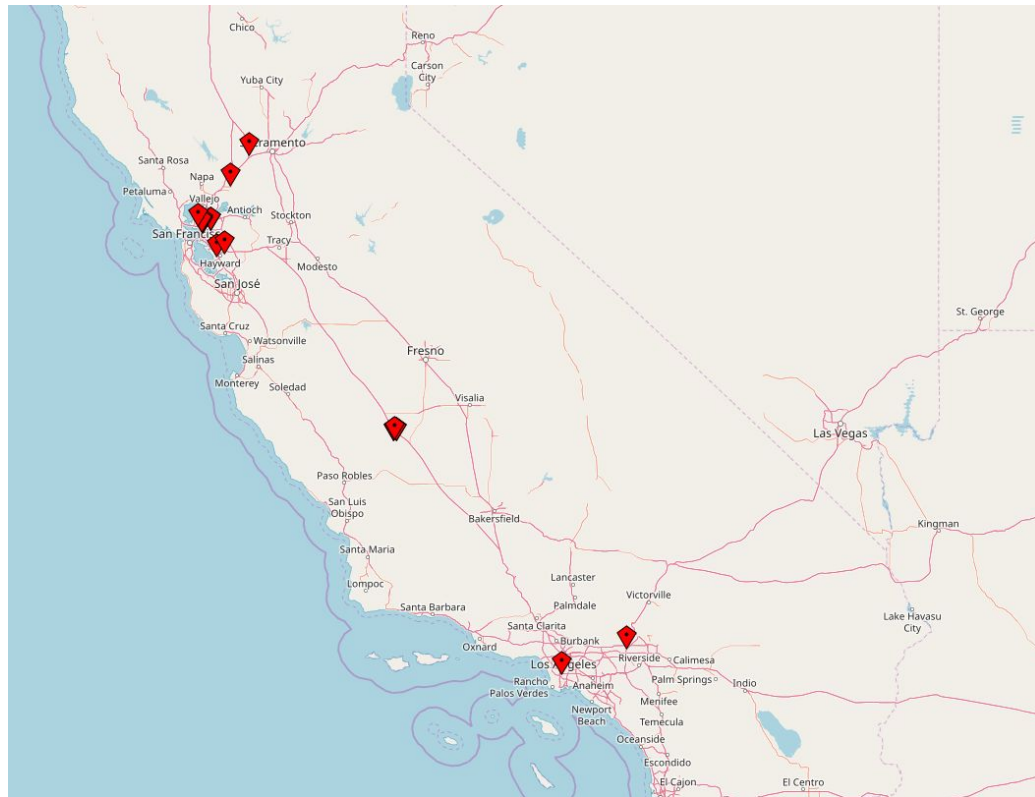
Dataset Overview

- **107** buildings
- **> 13.5 billion** data points
- **26,000** data streams
- All data streams described in per-site Brick models
- All data available on <https://mortardata.org/>

Temperature Sensor	7380
Luminance Sensor	257
Occupancy Sensor	445
Pressure Sensor	148
Outside Air Temperature Sensor	362
Cloud Cover	32
Setpoints (generic)	2331
Power Meters	77
VAVs	4724
AHUs	467
HVAC Zones	4887
Dampers	1662
Non BMS Thermostats	123

Dataset Overview

- UC Davis campus data
 - ~ 90 buildings (not all with data)
 - 3-6 years of data
 - ~15min collection interval
 - HVAC system, building meter
- XBOS-DR/V data
 - 15 RTU-based buildings
 - 1-9 months of data
 - ~10-30sec collection interval
 - HVAC, Lighting, building meter, occupancy, PV, EVSE



API for Heterogeneous Data

- How do apps actually make use of this data?
- Need data representations that are application-appropriate
 - Simplify expressive and semantically rich underlying data model
- Application-specific dataset as a **materialized view**
 - Expressed as queries over Brick metadata + timeseries data

Two Flavors of “Views”

Metadata (Context)

Timeseries Data



Two Flavors of “Views”

Metadata (Context)

```
temperature_sensors = pymortar.View(  
    name="temp_sensors",  
    definition="""SELECT ?temp ?room ?zone  
        WHERE {  
            ?temp rdf:type/rdfs:subClassOf*  
                brick:Temperature_Sensor .  
            ?temp bf:isLocatedIn ?room .  
            ?room rdf:type brick:Room .  
            ?room bf:isPartOf ?zone .  
            ?zone rdf:type brick:HVAC_Zone  
        };""",  
)
```

Declarative specification of app dataset

Timeseries Data

```
temperature_streams = pymortar.DataFrame(  
    name="temperature_sensor_data",  
    aggregation=pymortar.MEAN,  
    window="15m",  
    timeseries=[  
        pymortar.Timeseries(  
            view="temp_sensors",  
            dataVars=["?temp"],  
        ),  
    ]  
)
```



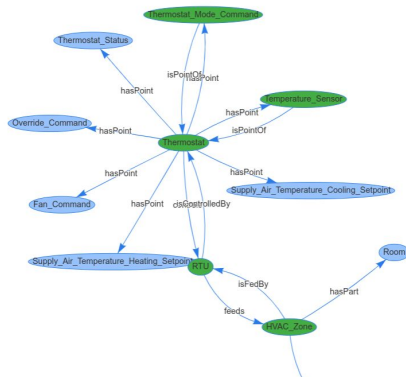
Metadata (Context)

Timeseries Data

Two Flavors of “Views”


Metadata (Context)

```
temperature_sensors = pymortar.View(
    name="temp_sensors",
    definition="""SELECT ?temp ?room ?zone
WHERE {
    ?temp rdf:type/rdfs:subClassOf*
        brick:Temperature_Sensor .
    ?temp bf:isLocatedIn ?room .
    ?room rdf:type brick:Room .
    ?room bf:isPartOf ?zone .
    ?zone rdf:type brick:HVAC_Zone
};""",
)
```



Timeseries Data

```
temperature_streams = pymortar.DataFrame(
    name="temperature_sensor_data",
    aggregation=pymortar.MEAN,
    window="15m",
    timeseries=[
        pymortar.Timeseries(
            view="temp_sensors",
            dataVars=["?temp"],
        ),
    ],
)
```



From SELECT clause			Autogenerated	
temp	room	zone	temp_uuid	site
sen1	Room 410	Zone 4	...-978b8df63cf0	CompSciHall
sen2	Room 410	Zone 4	...-d7c96010da5c	CompSciHall
sen3	Room 411	Zone 4	...-b1607a419aa4	CompSciHall
sen4	Room 501	Zone 5	...-89f84cf4ac34a	CompSciHall
...

Yields app-specific
(relational) table

	fbee4857-0a4e-3a18-808d-c94ad6801642	de32b7a8-38ae-322b-a98f-4d4dddd84e37	7d49226d-32a0-328f-b8cb-5c3adddbd40c5	dfde85ff-0a94-3e4f-9dc1-05921a8408f
2018-08-31 21:30:00+00:00	1306.0	52.0	105.0	Na
2018-08-31 22:00:00+00:00	1312.0	52.0	105.0	Na
2018-08-31 22:30:00+00:00	1312.0	52.0	105.0	Na
2018-08-31 23:00:00+00:00	1392.0	52.0	105.0	Na
2018-08-31 23:30:00+00:00	1350.0	52.0	105.0	Na



Two Flavors of “Views”



Metadata (Context)

```
temperature_sensors = pymortar.View(  
    name="temp_sensors",  
    definition="""SELECT ?temp ?room ?zone  
    WHERE {  
        ?temp rdf:type Temperature_Sensor  
        ?temp bf:hasPoint ?room  
        ?room rdf:type Room  
        ?room bf:hasZone ?zone  
        ?zone rdf:type Zone  
    }""",  
)
```

Timeseries Data

```
temperature_streams = pymortar.Timeseries(  
    name="temperature_streams",  
    aggregation=pymortar.AGGREGATION_MAX,  
    window="15m",  
    timeseries=[  
        pymortar.Timeseries(  
            view="temp_sensors",  
            dataVars=["?temp"],  
        ),  
    ],  
)
```

Declarative Mortar API allows applications to project heterogeneous data into a tractable form



From **SELECT** clause

Autogenerated

temp	room	zone	temp_uuid	site
n1	Room 410	Zone 4	...-9788bdf63ef0	CompSciHall
n2	Room 410	Zone 4	...-d7c96010da5c	CompSciHall
n3	Room 411	Zone 4	...-b1607a413aa4	CompSciHall
n4	Room 501	Zone 5	...-89f84ef4ac34a	CompSciHall
..


Yields app-specific
(relational) table

	f5ee4857-0a4e-3a18-808d-c94ad6801642	de32b7a8-38ae-322b-a98f-4d4ddddd84e37	7d49226d-32a0-328f-b8cb-5c3dddbd40c5	dfde85f0a94-3e4f9dc1-05921a8408f
2018-08-31 00:00	1306.0	52.0	105.0	Na
2018-08-31 22:00:00+00:00	1312.0	52.0	105.0	Na
2018-08-31 22:30:00+00:00	1312.0	52.0	105.0	Na
2018-08-31 23:00:00+00:00	1392.0	52.0	105.0	Na
2018-08-31 23:30:00+00:00	1350.0	52.0	105.0	Na



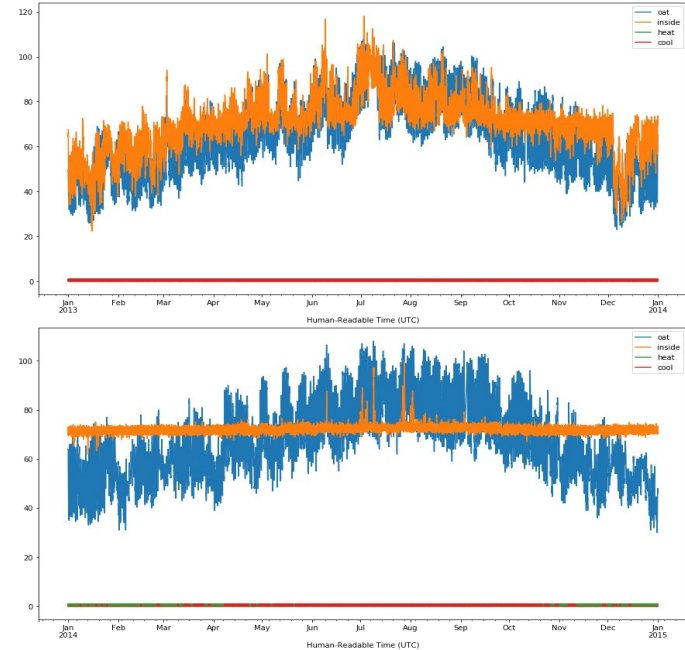
Example: Self-configuring Thermal Model

- Running optimal thermostatic control in ~20 buildings
 - Butcher shop, movie theatre, vets hall, senior center, office buildings, residential, etc
 - Different {equipment, sensors, architecture, climates}
- Requires data-driven thermal model:
 - ~~Hand code building specific thermal model~~
 - Code a thermal model *template* app
 - App uses Mortar to configure itself to a building
 - 20 different thermal models with the same code base



Looking Forward: Applying RL to Building Controls

- Reduce the role of the expert in designing control schemes for buildings
 - Expert doesn't scale
- Need robust simulations to apply RL
 - Work on thermal model enables creating “digital twins” at scale
- How to write/deploy RL agents (at scale) for different deployment context?



Outline

- ~~1. Lecture on Brick Schema~~
- ~~2. Overview of Mortar platform~~
3. Guided exploration of Mortar with online tutorial

Resources

Brick Website

<https://brickschema.org/>

Mortar Website

<https://mortardata.org/>

Tutorial Page

<https://tutorial.mortardata.org/>

Brick Explore Tool

<https://querybuilder.mortardata.org/>

Interactive Portion

1. Go to mortardata.org and click “Sign Up” to make an account
2. Go to tutorial.mortardata.org
3. Complete the first 2 cells